

Ledger as an auditing solution for cloud database By CloudChain

Team Members

Akash Malla

Sharon Subathran

Neha Soma

Sneha Mule





Outline

- Introduction
- Implementation
- Demo
- Conclusion
- Challenges
- Take Away



Introduction

Aim:

- To improve the security and consistency of data in a NoSQL database namely mongoDB.
- With further research, we felt that we could not contribute anymore than the existing solution (Crypt-NoSQL,).
- To provide a concise and clean platform to audit logs in a NoSQL database such as mongoDB.



Problem Statement

- Typically, many clients have to extract data from database logs to analyze performance issues.
- Reviewing and formatting log data is time consuming, if not, more time consuming than analyzing formatted data using a visualization utility or machine learning algorithm
- We wanted to eliminate the heavy lifting of formatting and capture all active operation related data in a single collection for any client to be able to query.

Why MongoDB?

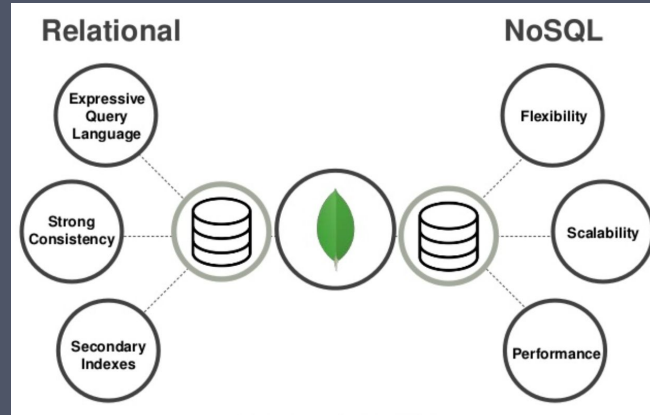
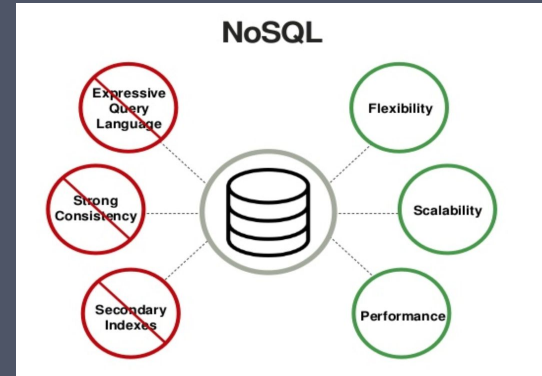
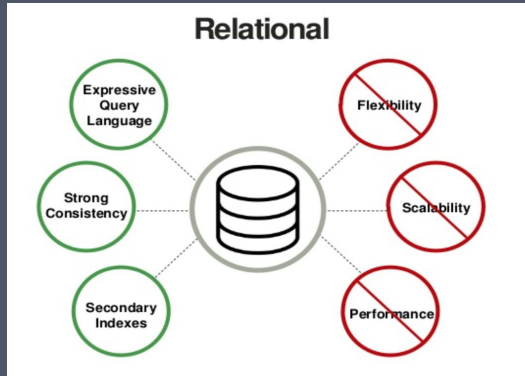


Image Ref : [*when-to-use-mongodb-when-you-should-not*](#)



Current MongoDB Tools

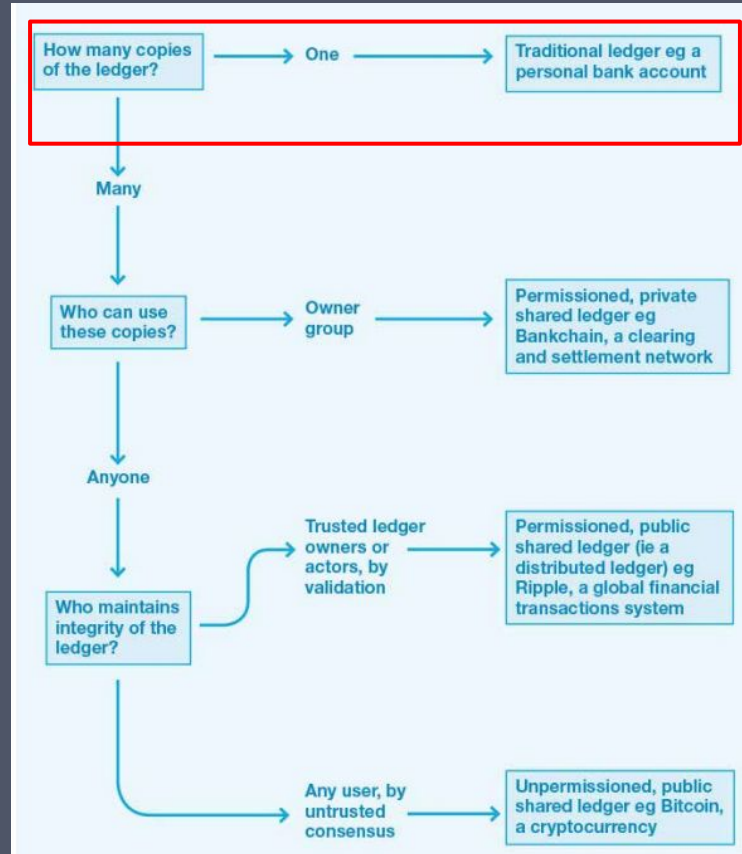
- MongoDB Atlas:
 - It takes data in JSON format.
 - It creates a cluster and gives visual representation.
- Coralogix:
 - It has a button for loggregator.
 - It gives pattern which collects same log clusters.
- AWS cloudWatch:
 - It can set alert based on threshold set.
 - It can store log data from multiple resource LogGroup



Distributed Ledger

- Distributed Ledger technology (DLT) may have started off as the basis for bitcoin, but it already promises to be much more than a cryptocurrency.
- Records/logs are stored one after the other in a continuous ledger, rather than sorted into blocks, but they can only be added when the participants reach a quorum.
- Currently, there are a number of digital ledgers that are emerging. Among them are ledgers which don't communicate with each other such as **Hyperledger Fabric, R3's Corda, Ripple, and more.**

Distributed Ledger Classification





Implementation Overview


- In our architecture, we built the ledger as a way to audit the mongodb logs instead of having large text files.
- This ledger is a collection, so it consists of objects which are easier to analyze and find valuable data for a consumer in order to determine what kind of queries are taking time.



Implementation Overview (contd..)

- The ledger facilitates a way to ensure that the right data is fetched and added into the ledger by using multiple mongoDB commands.
- List of operations in the ledger can be accessed easily from the mongodb querying platform by all clients.

Our Design



Ledger Parameters used:

- Operation id and type
- Microseconds Running
- Client IP address
- Namespace/Collection on which operation is running on
- Current and available connections
- Network byte input and byte output
- Network number of requests
- Total number of inserts, queries, updates at the moment
- Total number of updates at the moment
- Number of page faults
- Virtual and resident memory of a process

Image of our ledger:

```
_id: ObjectId("5aaaf698d49ad729c08f58db")
opid: 291619
operation: "update"
microsecs_running: 330
client: "127.0.0.1:36720"
namespace: "POCDB.POCCOLL5"
time: 2018-03-15 08:41:28.322
current_connections: 9
available_connections: 51191
active_clients: 16
current_queue: 0
network_bytesIn: 86200665
network_bytesOut: 56528302
network_numRequests: 291360
opcounters_insert: 144862
opcounters_query: 144414
opcounters_update: 143942
extra_info_page_faults: 193
memory_virtual: 1029
memory_resident: 738
```



Tools we used

- POCDriver.jar (<https://github.com/johnlpage/POCDriver>) - For load testing mongoDB instance.
- Amazon AWS EC2 (Ubuntu 16.04.3)
- Mongo Compass - For visualizing ledger data and mongodb data other than mongo shell
- Pymongo API in order to connect to mongodb in python
- Mongo shell

Implementation in Detail

Below is how we kick start the load test:

```
[ubuntu@ip-172-31-27-35:~]$ java -jar POCDriver.jar -k 20 -i 20 -u 20 -d 600 -t 3 -e -y 10
MongoDB Proof Of Concept - Load Generator
Worker thread 0 Started.
Worker thread 1 Started.
Worker thread 2 Started.
-----
After 10 seconds, 15165 new records inserted - collection has 15165 in total
1387 inserts per second since last report 0.00 % in under 50 milliseconds
144 keyqueries per second since last report 100.00 % in under 50 milliseconds
1453 updates per second since last report 90.48 % in under 50 milliseconds
0 rangequeries per second since last report 100.00 % in under 50 milliseconds
-----
After 20 seconds, 36645 new records inserted - collection has 36645 in total
2148 inserts per second since last report 0.00 % in under 50 milliseconds
210 keyqueries per second since last report 100.00 % in under 50 milliseconds
2157 updates per second since last report 90.36 % in under 50 milliseconds
0 rangequeries per second since last report 100.00 % in under 50 milliseconds
```

Mongostat to check performance

```
[ubuntu@ip-172-31-27-35:~$ mongostat
```

insert	query	update	delete	getmore	command	% dirty	% used	flushes	vsize	res	qr qw	ar aw	netIn	netOut	conn	time
*0	*0	*0	*0	0	14 0	0.0	0.1	0	280M	53.0M	0 0	0 0	1.99k	52.7k	4	2018-03-15T21:09:53Z
*0	*0	*0	*0	0	18 0	0.0	0.1	0	280M	53.0M	0 0	0 0	1.67k	286k	3	2018-03-15T21:09:54Z
*0	*0	*0	*0	0	19 0	0.0	0.1	0	280M	53.0M	0 0	0 0	1.74k	267k	2	2018-03-15T21:09:55Z
*0	*0	*0	*0	0	14 0	0.0	0.1	0	280M	53.0M	0 0	0 0	1.99k	52.1k	4	2018-03-15T21:09:56Z
339	416	331	*0	0	34 0	0.1	1.9	0	304M	73.0M	0 0	1 0	216k	343k	5	2018-03-15T21:09:57Z
1016	1100	1181	*0	0	26 0	0.5	5.2	0	339M	106M	0 0	1 0	636k	551k	7	2018-03-15T21:09:58Z
1429	1258	1493	*0	0	10 0	1.1	7.7	0	364M	130M	0 0	0 0	1.08m	389k	9	2018-03-15T21:09:59Z
864	1046	13980	*0	0	24 0	1.3	8.6	0	374M	140M	0 0	0 1	1.67m	592k	5	2018-03-15T21:10:00Z
1854	1709	1949	*0	0	10 0	2.0	10.3	0	391M	157M	0 0	0 1	1.10m	452k	7	2018-03-15T21:10:01Z
1621	1652	1667	*0	0	47 0	2.5	11.5	0	405M	170M	0 0	0 0	977k	1.11m	5	2018-03-15T21:10:02Z
insert	query	update	delete	getmore	command	% dirty	% used	flushes	vsize	res	qr qw	ar aw	netIn	netOut	conn	time
1896	1828	1793	*0	0	28 0	3.1	12.7	0	418M	182M	0 0	0 0	1.11m	692k	9	2018-03-15T21:10:03Z
1792	1822	1636	*0	0	24 0	3.7	13.9	0	430M	194M	0 0	0 0	1.05m	817k	7	2018-03-15T21:10:04Z
2032	1783	1997	*0	0	20 0	4.3	15.0	0	443M	206M	0 0	0 0	1.17m	585k	7	2018-03-15T21:10:05Z
1727	2079	1809	*0	0	18 0	5.0	16.2	0	456M	218M	0 0	0 0	1.08m	744k	8	2018-03-15T21:10:06Z
2029	1815	2035	*0	0	10 0	5.2	17.4	0	468M	229M	0 0	0 0	1.18m	459k	5	2018-03-15T21:10:07Z
2080	2143	2043	*0	0	27 0	5.3	18.6	0	482M	242M	0 0	0 0	1.24m	822k	5	2018-03-15T21:10:08Z
2220	2159	2177	*0	0	28 0	5.3	19.9	0	495M	256M	0 0	0 0	1.31m	861k	5	2018-03-15T21:10:09Z
2185	2003	2160	*0	0	10 0	4.9	21.0	0	508M	268M	0 0	0 0	1.28m	503k	8	2018-03-15T21:10:10Z
2090	2282	2148	*0	0	24 0	5.1	22.3	0	522M	281M	0 0	1 0	1.28m	877k	5	2018-03-15T21:10:11Z
2063	1924	2063	*0	0	6 0	4.8	23.4	0	535M	293M	0 0	1 0	1.21m	480k	6	2018-03-15T21:10:12Z
insert	query	update	delete	getmore	command	% dirty	% used	flushes	vsize	res	qr qw	ar aw	netIn	netOut	conn	time
2050	2102	2067	*0	0	30 0	4.8	24.6	0	548M	304M	0 0	0 0	1.23m	826k	5	2018-03-15T21:10:13Z
2049	2059	2007	*0	0	16 0	4.7	25.7	0	561M	316M	0 0	0 0	1.22m	548k	8	2018-03-15T21:10:14Z
2142	2141	2173	*0	0	21 0	4.7	26.8	0	573M	327M	0 0	0 0	1.28m	706k	7	2018-03-15T21:10:15Z
2425	2218	2210	*0	0	18 0	5.2	28.2	0	588M	341M	0 0	0 0	1.39m	658k	5	2018-03-15T21:10:16Z
1994	2165	2052	*0	0	32 0	5.2	29.2	0	600M	352M	0 0	1 0	1.22m	832k	5	2018-03-15T21:10:17Z
2723	2444	2723	*0	0	22 0	6.2	30.6	0	614M	364M	0 0	0 0	1.59m	846k	9	2018-03-15T21:10:18Z
2403	2460	2451	*0	0	15 0	5.7	31.9	0	627M	376M	0 0	0 0	1.44m	656k	5	2018-03-15T21:10:19Z
2651	2668	2548	*0	0	17 0	5.6	33.2	0	642M	389M	0 0	0 0	1.57m	661k	11	2018-03-15T21:10:20Z
2629	2664	2664	*0	0	23 0	5.5	34.6	0	657M	402M	0 0	0 0	1.57m	1.03m	5	2018-03-15T21:10:21Z
2690	2701	2720	*0	0	31 0	5.3	35.9	0	672M	416M	0 0	1 1	1.60m	936k	7	2018-03-15T21:10:22Z

Code

- `mongodb_conn()` function tries to connect to mongodb server in 10 seconds.

```
import pymongo
import schedule
import time
import datetime
import traceback

def mongodb_conn():
    try:
        maxSevSelDelay=10000
        client = pymongo.MongoClient("13.56.82.17",serverSelectionTimeoutMS=maxSevSelDelay)
        client.server_info()
    except pymongo.errors.ConnectionFailure as err1:
        print("Could not connect to server, connection failure:",err1)
    except pymongo.errors.ServerSelectionTimeoutError as err2:
        print("Could not connect to server, ServerSelectionTimeoutError:",err2)
    return client
```



```

def job():
    conn = mongodb_conn()
    #print("connected successfully")
    if conn is None:
        # no connection, exit early
        print("not connected to mongodb!")
        return

    #print("I'm working...")
    try:
        db = conn.data
        input=db.current_op(True)

        for j in input['inprog']:

            if j['active']==True:
                if 'microsecs_running' in j and j['microsecs_running']>50 and j['microsecs_running']<15000:
                    #print(j)
                    newJ=dict()
                    if 'opid' in j:
                        newJ['opid']=j['opid']
                    if 'op' in j:
                        newJ['operation']=j['op']
                    if 'microsecs_running' in j:
                        newJ['microsecs_running']=j['microsecs_running']
                    if 'client' in j:
                        newJ['client']=j['client']
                    if 'ns' in j:
                        newJ['namespace']=j['ns']
                    newJ['time']=datetime.datetime.now()
                    server_status=db.command({"serverStatus":1})
                    newJ['current_connections']=server_status['connections']['current']
                    newJ['available_connections']=server_status['connections']['available']
                    newJ['active_clients']=server_status['globalLock']['activeClients']['total']
                    newJ['current_queue']=server_status['globalLock']['currentQueue']['total']
                    newJ['network_bytesIn']=server_status['network']['bytesIn']
                    newJ['network_bytesOut']=server_status['network']['bytesOut']
                    newJ['network_numRequests']=server_status['network']['numRequests']
                    newJ['opcounters_insert']=server_status['opcounters']['insert']
                    newJ['opcounters_query']=server_status['opcounters']['query']
                    newJ['opcounters_update']=server_status['opcounters']['update']

```

```

newJ['network_numRequests']=server_status['network']['numRequests']
newJ['opcounters_insert']=server_status['opcounters']['insert']
newJ['opcounters_query']=server_status['opcounters']['query']
newJ['opcounters_update']=server_status['opcounters']['update']
newJ['extra_info_page_faults']=server_status['extra_info']['page_faults']
newJ['memory_virtual']=server_status['mem']['virtual']
newJ['memory_resident']=server_status['mem']['resident']
print(newJ)
#print("inside\n")

db.ledgerData.insert_one(newJ)

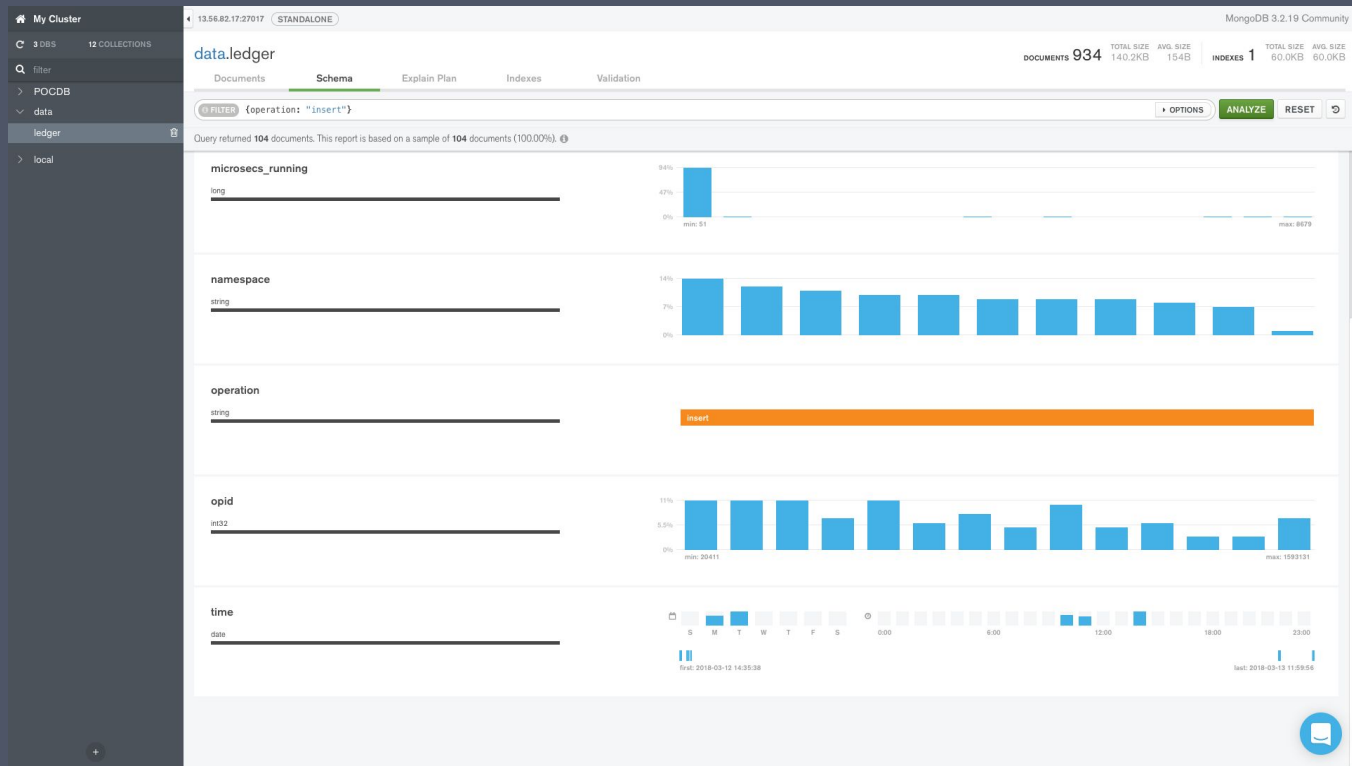
except:
    traceback.print_exc()
    exit(1)

job()
schedule.every(1).seconds.do(job)
schedule.every().hour.do(job)
schedule.every().day.do(job)

while 1:
    schedule.run_pending()
    time.sleep(1)

```


Demo






Is it advisable to implement our solution?

- This solution is advisable assuming space is not a constraint for a client. In a cloud environment, space constraint can be resolved by simply adding more machines (horizontally Scaling).
- Therefore, as a short answer: Yes, it is advisable, however, if adding more machines leads to going over a client's budget, then this solution may not be advisable in that scenario.



Where would it work well and where would it not (tradeoffs)?

- Main trade off will be a memory consumption.
- It is a concise version of typical NoSQL log entries.
- It will work well while performing analysis of the data so that service provider can take action in advance to prevent query delay.



Who would benefit from such a solution and what are the possible areas where it can be applied?

- Ability to monitor data in heavy traffic time (For example: Thanksgiving weekend) and Low Traffic time during the course of an year.
- Customer can choose an appropriate time interval to insert into ledger (For example: every second, 30 seconds, a minute or more).
- Customer can run machine learning algorithms to analyze the performance of each operation type.
- Customer has flexibility to keep or remove ledger data depending on how much space available in their cloud environment.



Problems Faced

- Installing mongoDB on **docker container** .
- Considered using **Mongo Perf**, however, it was not well documented and properly tested as we kept facing errors running the utility.
- There was no documentation or video tutorials for stress testing using **Jmeter for mongoDB**.
- Understanding blockchain technology was time consuming and implementing it in our solution.



Future Work

- Use this data to do analysis using machine learning algorithm to predict or detect query which are running slowly and take precautionary action in advance.
- Add visualization like clustering data , error detection , graphical representation.



Our Takeaways

- Concepts of Blockchain Technology :
 - Ledger
 - HashCode
 - SHA-256
- How Blockchain used in cryptocurrency
- Working with Jmeter, Mongo Perf
- Working with MongoDB and its commands (mongostat, db.server_status(), mongotop etc)
- Working with the above using Amazon EC2 instance



References

1. “*Building Enterprise-Grade Blockchain Databases with MongoDB*” white paper, 2017.
<https://www.mongodb.com/white-papers>
2. A. Stanciu, "*Blockchain Based Distributed Control System for Edge Computing*," 2017 21st International Conference on Control Systems and Computer Science (CSCS), Bucharest, 2017, pp. 667-671.
3. “MongoDB Documentaion ” <https://docs.mongodb.com/manual/tutorial/rotate-log-files/>
4. “MongoDB Compass” <https://www.mongodb.com/products/compass>
5. “PyMongo Documentation” <https://api.mongodb.com/python/current/>
6. “Apache Jmeter” <https://jmeter.apache.org/>
7. <https://docs.aws.amazon.com/quickstart/latest/mongodb/architecture.html>



Questions?



Thank you!